

STA9750__Project__Code

Soonmo Seong

12/15/2019

Project Code

```
library(dplyr)
library(MASS)
library(glmnet)
library(ggplot2)
library(reshape)
library(gridExtra)
library(xtable)

a = read.csv("pa.csv")
dim(a)
t.data = a[,-(1:2)]

# standardize the dataset
t.data[,-(1:2)] = scale(t.data[,-(1:2)])
t.data = t.data[, -2]
colMeans(t.data)
apply(t.data, 2, sd)
mean(t.data$Status)
attach(t.data)

# data structure and rates
n1 = dim(t.data %>% filter(Status == 1))[1]
n0 = dim(t.data %>% filter(Status == 0))[1]
n = n1 + n0
p = dim(t.data)[2]-1

#####
#####
# Modelling factors
iterations = 100
Dlearn_rate = 0.8
sampling.rate = 1

# train, test, and cv error rate matrix
train_error = matrix(0, nrow = iterations, ncol = 3)
colnames(train_error) = c("Log", "Log-LASSO", "Log-Ridge")

cv_error = matrix(0, nrow = iterations, ncol = 2)
colnames(cv_error) = c("Log-LASSO", "Log-Ridge")

test_error = matrix(0, nrow = iterations, ncol = 3)
colnames(test_error) = c("Log", "Log-LASSO", "Log-Ridge")

cv_error_nfold = matrix(0, nrow = iterations, ncol = 30)
```

```

colnames(cv_error_nfold) = c("3_Log-LASSO",
                             "4_Log-LASSO", "5_Log-LASSO",
                             "6_Log-LASSO", "7_Log-LASSO",
                             "8_Log-LASSO", "9_Log-LASSO",
                             "10_Log-LASSO",
                             "15_Log-LASSO", "20_Log-LASSO",
                             "25_Log-LASSO", "30_Log-LASSO",
                             "40_Log-LASSO", "50_Log-LASSO",
                             "L00_Log-LASSO",
                             "3_Log-Ridge",
                             "4_Log-Ridge", "5_Log-Ridge",
                             "6_Log-Ridge", "7_Log-Ridge",
                             "8_Log-Ridge", "9_Log-Ridge",
                             "10_Log-Ridge",
                             "15_Log-Ridge", "20_Log-Ridge",
                             "25_Log-Ridge", "30_Log-Ridge",
                             "40_Log-Ridge", "50_Log-Ridge",
                             "L00_Log-Ridge")

lasso.coef = matrix(0, ncol = iterations, nrow = p+1)
ridge.coef = matrix(0, ncol = iterations, nrow = p+1)

# convert to data frame
train_error = data.frame(train_error)
test_error = data.frame(test_error)
cv_error = data.frame(cv_error)
cv_error_nfold = data.frame(cv_error_nfold)

# time of cv and fit
time.cv = matrix(0, nrow = iterations, ncol = 2)
colnames(time.cv) = c("LASSO", "Ridge")

time.cv.nfold = matrix(0, nrow = iterations, ncol = 30)
colnames(time.cv.nfold) = c("3_Log-LASSO",
                             "4_Log-LASSO", "5_Log-LASSO",
                             "6_Log-LASSO", "7_Log-LASSO",
                             "8_Log-LASSO", "9_Log-LASSO",
                             "10_Log-LASSO",
                             "15_Log-LASSO", "20_Log-LASSO",
                             "25_Log-LASSO", "30_Log-LASSO",
                             "40_Log-LASSO", "50_Log-LASSO",
                             "L00_Log-LASSO",
                             "3_Log-Ridge",
                             "4_Log-Ridge", "5_Log-Ridge",
                             "6_Log-Ridge", "7_Log-Ridge",
                             "8_Log-Ridge", "9_Log-Ridge",
                             "10_Log-Ridge",
                             "15_Log-Ridge", "20_Log-Ridge",
                             "25_Log-Ridge", "30_Log-Ridge",
                             "40_Log-Ridge", "50_Log-Ridge",
                             "L00_Log-Ridge")

time.fit = matrix(0, nrow = iterations, ncol = 3)

```

```

colnames(time.fit) = c("Log", "LASSO", "Ridge")

# sampling from t.data(ntrain, ntest)
sampling = sample(n,n*sampling.rate)
sampling.data = data.frame(t.data[sampling,])
sampling.n = dim(sampling.data)[1]

# preparation for lasso and ridge
X = model.matrix(Status ~., sampling.data)[,-1]
y = sampling.data$Status

# 100 iteration for training error rate and testing error rate
for(m in 41:iterations){

  # create a training data vector for dividing the data set.
  train = sample(sampling.n, sampling.n*Dlearn_rate)

  dat  = data.frame(sampling.data[train,])
  datt = data.frame(sampling.data[-train,])

  # logistic regression
  log.mod = glm(Status ~., data = dat, family = "binomial")
  log.pred = predict(log.mod, newdata = sampling.data[train,], type = "response")
  log.pred = ifelse(log.pred > 0.5, 1, 0)
  train_error[m,1] = mean( sampling.data[train,1] != log.pred)
  log.pred = predict(log.mod, newdata = sampling.data[-train,], type = "response")
  log.pred = ifelse(log.pred > 0.5, 1, 0)
  test_error[m,1] = mean( sampling.data[-train,1] != log.pred)

  # lasso cross validation and tune lambda
  # record lasso cv time and error rate
  ptm = proc.time()
  cv.lasso = cv.glmnet(X[train,], y[train], alpha = 1, family = "binomial",
                      intercept = T, type.measure="class")
  ptm = proc.time() - ptm
  time.cv[m,1] = ptm["elapsed"]

  cv_error[m,1] = min(cv.lasso$cvm)
  bestlam = cv.lasso$lambda.min
  # record lasso fit time
  ptm = proc.time()
  lasso.mod = glmnet(X[train,], y[train], alpha = 1, family = "binomial",
                    intercept = T, lambda = bestlam,
                    standardize = F)
  ptm = proc.time() - ptm
  time.fit[m,2] = ptm["elapsed"]

  lasso.coef[,m] = coef(lasso.mod)[,1]
  lasso.pred = predict(lasso.mod, s = bestlam, newx = X[train,], type ="class")
  train_error[m,2] = mean(y[train] != lasso.pred)
  lasso.pred = predict(lasso.mod, s = bestlam, newx = X[-train,], type="class",)

```

```

test_error[m,2] = mean(y[-train] != lasso.pred)

# record lasso cv time and error rate by nfolds
# 3fold cv lasso
ptm = proc.time()
cv.lasso = cv.glmnet(X[train,], y[train], alpha = 1, family = "binomial",
                    intercept = T, type.measure="class", nfolds = 3)
ptm = proc.time() - ptm
time.cv.nfold[m,1] = ptm["elapsed"]

cv_error_nfold[m,1] = min(cv.lasso$cvm)
# 4fold cv lasso
ptm = proc.time()
cv.lasso = cv.glmnet(X[train,], y[train], alpha = 1, family = "binomial",
                    intercept = T, type.measure="class", nfolds = 4)
ptm = proc.time() - ptm
time.cv.nfold[m,2] = ptm["elapsed"]

cv_error_nfold[m,2] = min(cv.lasso$cvm)
# 5fold cv lasso
ptm = proc.time()
cv.lasso = cv.glmnet(X[train,], y[train], alpha = 1, family = "binomial",
                    intercept = T, type.measure="class", nfolds = 5)
ptm = proc.time() - ptm
time.cv.nfold[m,3] = ptm["elapsed"]

cv_error_nfold[m,3] = min(cv.lasso$cvm)
# 6fold cv lasso
ptm = proc.time()
cv.lasso = cv.glmnet(X[train,], y[train], alpha = 1, family = "binomial",
                    intercept = T, type.measure="class", nfolds = 6)
ptm = proc.time() - ptm
time.cv.nfold[m,4] = ptm["elapsed"]

cv_error_nfold[m,4] = min(cv.lasso$cvm)
# 7fold cv lasso
ptm = proc.time()
cv.lasso = cv.glmnet(X[train,], y[train], alpha = 1, family = "binomial",
                    intercept = T, type.measure="class", nfolds = 7)
ptm = proc.time() - ptm
time.cv.nfold[m,5] = ptm["elapsed"]

cv_error_nfold[m,5] = min(cv.lasso$cvm)
# 8fold cv lasso
ptm = proc.time()
cv.lasso = cv.glmnet(X[train,], y[train], alpha = 1, family = "binomial",
                    intercept = T, type.measure="class", nfolds = 8)
ptm = proc.time() - ptm
time.cv.nfold[m,6] = ptm["elapsed"]

cv_error_nfold[m,6] = min(cv.lasso$cvm)
# 9fold cv lasso
ptm = proc.time()

```

```

cv.lasso = cv.glmnet(X[train,], y[train], alpha = 1, family = "binomial",
                    intercept = T, type.measure="class", nfolds = 9)
ptm = proc.time() - ptm
time.cv.nfold[m,7] = ptm["elapsed"]

cv_error_nfold[m,7] = min(cv.lasso$cvm)
# 10fold cv lasso
ptm = proc.time()
cv.lasso = cv.glmnet(X[train,], y[train], alpha = 1, family = "binomial",
                    intercept = T, type.measure="class", nfolds = 10)
ptm = proc.time() - ptm
time.cv.nfold[m,8] = ptm["elapsed"]

cv_error_nfold[m,8] = min(cv.lasso$cvm)
# 15fold cv lasso
ptm = proc.time()
cv.lasso = cv.glmnet(X[train,], y[train], alpha = 1, family = "binomial",
                    intercept = T, type.measure="class", nfolds = 15)
ptm = proc.time() - ptm
time.cv.nfold[m,9] = ptm["elapsed"]

cv_error_nfold[m,9] = min(cv.lasso$cvm)
# 20fold cv lasso
ptm = proc.time()
cv.lasso = cv.glmnet(X[train,], y[train], alpha = 1, family = "binomial",
                    intercept = T, type.measure="class", nfolds = 20)
ptm = proc.time() - ptm
time.cv.nfold[m,10] = ptm["elapsed"]

cv_error_nfold[m,10] = min(cv.lasso$cvm)
# 25fold cv lasso
ptm = proc.time()
cv.lasso = cv.glmnet(X[train,], y[train], alpha = 1, family = "binomial",
                    intercept = T, type.measure="class", nfolds = 25)
ptm = proc.time() - ptm
time.cv.nfold[m,11] = ptm["elapsed"]

cv_error_nfold[m,11] = min(cv.lasso$cvm)
# 30fold cv lasso
ptm = proc.time()
cv.lasso = cv.glmnet(X[train,], y[train], alpha = 1, family = "binomial",
                    intercept = T, type.measure="class", nfolds = 30)
ptm = proc.time() - ptm
time.cv.nfold[m,12] = ptm["elapsed"]

cv_error_nfold[m,12] = min(cv.lasso$cvm)
# 40fold cv lasso
ptm = proc.time()
cv.lasso = cv.glmnet(X[train,], y[train], alpha = 1, family = "binomial",
                    intercept = T, type.measure="class", nfolds = 40)
ptm = proc.time() - ptm
time.cv.nfold[m,13] = ptm["elapsed"]

```

```

cv_error_nfold[m,13] = min(cv.lasso$cvm)
# 50fold cv lasso
ptm = proc.time()
cv.lasso = cv.glmnet(X[train,], y[train], alpha = 1, family = "binomial",
                    intercept = T, type.measure="class", nfolds = 50)
ptm = proc.time() - ptm
time.cv.nfold[m,14] = ptm["elapsed"]
cv_error_nfold[m,14] = min(cv.lasso$cvm)
# LOO cv lasso
ptm = proc.time()
cv.lasso = cv.glmnet(X[train,], y[train], alpha = 1, family = "binomial",
                    intercept = T, type.measure="class", nfolds = length(train))
ptm = proc.time() - ptm
time.cv.nfold[m,15] = ptm["elapsed"]

cv_error_nfold[m,15] = min(cv.lasso$cvm)

# ridge cross validation and tune lambda
# record ridge cv time
ptm = proc.time()
cv.ridge = cv.glmnet(X[train,], y[train], alpha = 0, family = "binomial",
                    intercept = T, type.measure="class")
ptm = proc.time() - ptm
time.cv[m,2] = ptm["elapsed"]

cv_error[m,2] = min(cv.ridge$cvm)
bestlam = cv.ridge$lambda.min

# record ridge fit time
ptm = proc.time()
ridge.mod = glmnet(X[train,], y[train], alpha = 0, family = "binomial",
                  intercept = T, lambda = bestlam,
                  standardize = F)
ptm = proc.time() - ptm
time.fit[m,3] = ptm["elapsed"]

ridge.coef[,m] = as.matrix(coef(ridge.mod))
ridge.pred = predict(ridge.mod, s = bestlam, newx = X[train,], type = "class")
train_error[m,3] = mean(y[train] != ridge.pred)
ridge.pred = predict(ridge.mod, s = bestlam, newx = X[-train,], type = "class")
test_error[m,3] = mean(y[-train] != ridge.pred)

# record ridge cv time and error rate by nfolds
# 3fold cv ridge
ptm = proc.time()
cv.ridge = cv.glmnet(X[train,], y[train], alpha = 0, family = "binomial",
                    intercept = T, type.measure="class", nfolds = 3)
ptm = proc.time() - ptm
time.cv.nfold[m,16] = ptm["elapsed"]

cv_error_nfold[m,16] = min(cv.ridge$cvm)
# 4fold cv ridge
ptm = proc.time()

```

```

cv.ridge = cv.glmnet(X[train,], y[train], alpha = 0, family = "binomial",
                    intercept = T, type.measure="class", nfolds = 4)
ptm = proc.time() - ptm
time.cv.nfold[m,17] = ptm["elapsed"]

cv_error_nfold[m,17] = min(cv.ridge$cvm)
# 5fold cv ridge
ptm = proc.time()
cv.ridge = cv.glmnet(X[train,], y[train], alpha = 0, family = "binomial",
                    intercept = T, type.measure="class", nfolds = 5)
ptm = proc.time() - ptm
time.cv.nfold[m,18] = ptm["elapsed"]

cv_error_nfold[m,18] = min(cv.ridge$cvm)
# 6fold cv ridge
ptm = proc.time()
cv.ridge = cv.glmnet(X[train,], y[train], alpha = 0, family = "binomial",
                    intercept = T, type.measure="class", nfolds = 6)
ptm = proc.time() - ptm
time.cv.nfold[m,19] = ptm["elapsed"]

cv_error_nfold[m,19] = min(cv.ridge$cvm)
# 7fold cv ridge
ptm = proc.time()
cv.ridge = cv.glmnet(X[train,], y[train], alpha = 0, family = "binomial",
                    intercept = T, type.measure="class", nfolds = 7)
ptm = proc.time() - ptm
time.cv.nfold[m,20] = ptm["elapsed"]

cv_error_nfold[m,20] = min(cv.ridge$cvm)
# 8fold cv ridge
ptm = proc.time()
cv.ridge = cv.glmnet(X[train,], y[train], alpha = 0, family = "binomial",
                    intercept = T, type.measure="class", nfolds = 8)
ptm = proc.time() - ptm
time.cv.nfold[m,21] = ptm["elapsed"]

cv_error_nfold[m,21] = min(cv.ridge$cvm)
# 9fold cv ridge
ptm = proc.time()
cv.ridge = cv.glmnet(X[train,], y[train], alpha = 0, family = "binomial",
                    intercept = T, type.measure="class", nfolds = 9)
ptm = proc.time() - ptm
time.cv.nfold[m,22] = ptm["elapsed"]

cv_error_nfold[m,22] = min(cv.ridge$cvm)
# 10fold cv ridge
ptm = proc.time()
cv.ridge = cv.glmnet(X[train,], y[train], alpha = 0, family = "binomial",
                    intercept = T, type.measure="class", nfolds = 10)
ptm = proc.time() - ptm
time.cv.nfold[m,23] = ptm["elapsed"]

```

```

cv_error_nfold[m,23] = min(cv.ridge$cvm)
# 15fold cv ridge
ptm = proc.time()
cv.ridge = cv.glmnet(X[train,], y[train], alpha = 0, family = "binomial",
                    intercept = T, type.measure="class", nfolds = 15)
ptm = proc.time() - ptm
time.cv.nfold[m,24] = ptm["elapsed"]

cv_error_nfold[m,24] = min(cv.ridge$cvm)
# 20fold cv ridge
ptm = proc.time()
cv.ridge = cv.glmnet(X[train,], y[train], alpha = 0, family = "binomial",
                    intercept = T, type.measure="class", nfolds = 20)
ptm = proc.time() - ptm
time.cv.nfold[m,25] = ptm["elapsed"]

cv_error_nfold[m,25] = min(cv.ridge$cvm)
# 25fold cv ridge
ptm = proc.time()
cv.ridge = cv.glmnet(X[train,], y[train], alpha = 0, family = "binomial",
                    intercept = T, type.measure="class", nfolds = 25)
ptm = proc.time() - ptm
time.cv.nfold[m,26] = ptm["elapsed"]

cv_error_nfold[m,26] = min(cv.ridge$cvm)
# 30fold cv ridge
ptm = proc.time()
cv.ridge = cv.glmnet(X[train,], y[train], alpha = 0, family = "binomial",
                    intercept = T, type.measure="class", nfolds = 30)
ptm = proc.time() - ptm
time.cv.nfold[m,27] = ptm["elapsed"]

cv_error_nfold[m,27] = min(cv.ridge$cvm)
# 40fold cv ridge
ptm = proc.time()
cv.ridge = cv.glmnet(X[train,], y[train], alpha = 0, family = "binomial",
                    intercept = T, type.measure="class", nfolds = 40)
ptm = proc.time() - ptm
time.cv.nfold[m,28] = ptm["elapsed"]

cv_error_nfold[m,28] = min(cv.ridge$cvm)
# 50fold cv ridge
ptm = proc.time()
cv.ridge = cv.glmnet(X[train,], y[train], alpha = 0, family = "binomial",
                    intercept = T, type.measure="class", nfolds = 50)
ptm = proc.time() - ptm
time.cv.nfold[m,29] = ptm["elapsed"]

cv_error_nfold[m,29] = min(cv.ridge$cvm)
# LOO cv ridge
ptm = proc.time()
cv.ridge = cv.glmnet(X[train,], y[train], alpha = 0, family = "binomial",
                    intercept = T, type.measure="class", nfolds = length(train))

```



```

ptm = proc.time() - ptm
time.cv.nfold[m,30] = ptm["elapsed"]

cv_error_nfold[m,30] = min(cv.ridge$cvm)
}

# tune lambda for lasso and ridge with optimal k-fold cross validation
# create a training data vector for dividing the data set.
train = sample(sampling.n, sampling.n*Dlearn_rate)

dat = data.frame(sampling.data[train,])
datt = data.frame(sampling.data[-train,])

# lasso 25 fold cross validation and tune lambda
cv.lasso = cv.glmnet(X[train,], y[train], alpha = 1, family = "binomial",
                    intercept = T, type.measure="class", nfolds = 25)
# ridge 10 fold cross validation and tune lambda
cv.ridge = cv.glmnet(X[train,], y[train], alpha = 0, family = "binomial",
                    intercept = T, type.measure="class", nfolds = 10)
lasso.cv.error = data.frame(lambda = cv.lasso$lambda, error = cv.lasso$cvm)
ridge.cv.error = data.frame(lambda = cv.ridge$lambda, error = cv.ridge$cvm)

# train and test error with best lambda
tune_train_error = matrix(0, nrow = iterations, ncol = 3)
colnames(tune_train_error) = c("Log", "Log-LASSO", "Log-Ridge")

tune_test_error = matrix(0, nrow = iterations, ncol = 3)
colnames(tune_test_error) = c("Log", "Log-LASSO", "Log-Ridge")

tune_train_error = data.frame(train_error)
tune_test_error = data.frame(test_error)

tune_lasso.coef = matrix(0, nrow = iterations, ncol = p+1)
tune_ridge.coef = matrix(0, nrow = iterations, ncol = p+1)

# 100 iteration with best lambdas
for(m in 1:iterations){

  # create a training data vector for dividing the data set.
  train = sample(sampling.n, sampling.n*Dlearn_rate)

  dat = data.frame(sampling.data[train,])
  datt = data.frame(sampling.data[-train,])

  # logistic regression
  log.mod = glm(Status ~., data = dat, family = "binomial")
  log.pred = predict(log.mod, newdata = sampling.data[train,], type = "response")
  log.pred = ifelse(log.pred > 0.5, 1, 0)
  tune_train_error[m,1] = mean( sampling.data[train,1] != log.pred)
  log.pred = predict(log.mod, newdata = sampling.data[-train,], type = "response")
  log.pred = ifelse(log.pred > 0.5, 1, 0)
  tune_test_error[m,1] = mean( sampling.data[-train,1] != log.pred)
}

```

```

# lasso regression
bestlam = cv.lasso$lambda.min
lasso.mod = glmnet(X[train,], y[train], alpha = 1, family = "binomial",
                  intercept = T, lambda = bestlam,
                  standardize = F)
tune_lasso.coef[m,] = coef(lasso.mod)[,1]
lasso.pred = predict(lasso.mod, s = bestlam, newx = X[train,], type = "class")
tune_train_error[m,2] = mean(y[train] != lasso.pred)
lasso.pred = predict(lasso.mod, s = bestlam, newx = X[-train,], type = "class",)
tune_test_error[m,2] = mean(y[-train] != lasso.pred)

# ridge regression
bestlam = cv.ridge$lambda.min
ridge.mod = glmnet(X[train,], y[train], alpha = 0, family = "binomial",
                  intercept = T, lambda = bestlam,
                  standardize = F)
tune_ridge.coef[m,] = coef(ridge.mod)[,1]
ridge.pred = predict(ridge.mod, s = bestlam, newx = X[train,], type = "class")
tune_train_error[m,3] = mean(y[train] != ridge.pred)
ridge.pred = predict(ridge.mod, s = bestlam, newx = X[-train,], type = "class")
tune_test_error[m,3] = mean(y[-train] != ridge.pred)
}

# best model fit
bestlam = cv.ridge$lambda.min
ridge.mod = glmnet(X, y, alpha = 0, family = "binomial",
                  intercept = T, lambda = bestlam,
                  standardize = F)
ridge.pred = predict(ridge.mod, s = bestlam, newx = X, type = "class")
mean(y != ridge.pred)
ridge.confusion = data.frame(table(ridge.pred, y))
xtable(ridge.confusion)

#####
#Graph Drawing
#####

#general comparison between logistic, lasso, and ridge regression
f1_1 = ggplot(melt(train_error), aes(x = variable, y = value, color = variable)) +
  geom_boxplot() + ylim(0,.5) + labs(x = element_blank(), y = "Misclassification Error Rate",
                                   title = "Train Error Rate") +
  scale_color_brewer(palette="Dark2") + theme(legend.position="none")
f1_2 = ggplot(melt(test_error), aes(x = variable, y = value, color = variable)) +
  geom_boxplot() + ylim(0,.5) + labs(x = element_blank(), y = "Misclassification Error Rate",
                                   title = "Test Error Rate") +
  scale_color_brewer(palette="Dark2") + theme(legend.position="none")
f1_3 = ggplot(melt(cv_error), aes(x = variable, y = value, color = variable)) +
  geom_boxplot() + ylim(0,.5) + labs(x = element_blank(), y = "Misclassification Error Rate",
                                   title = "CV Error Rate") +
  scale_color_brewer(palette="Dark2") + theme(legend.position="none")
f1 = grid.arrange(f1_1, f1_2, f1_3, nrow = 1, widths = c(3,3,2))

```

```

# cross validation folds for lasso
lasso.average.cv.time.and.error = data.frame(error = round(colMeans(cv_error_nfold[,1:15]),4),
      time = round(colMeans(time.cv.nfold[,1:15]),2),
      fold = factor(c("3","4","5","6","7",
        "8","9","10","15","20",
        "25","30","40","50","L00"),
        levels = c("3","4","5","6","7",
          "8","9","10","15","20",
          "25","30","40","50","L00")))

f2_1 = ggplot() +
  geom_line(data = lasso.average.cv.time.and.error, aes(x = fold, y = error, group = 1,
    color = "line")) + ylim(0.15,.2) +
  geom_point(data = lasso.average.cv.time.and.error, aes(x = fold[which.min(error)],
    y = min(error), color = "min")) +
  labs(x = "Cross Validation Folds", y = "Misclassification Error Rate",
    title = "LASSO CV Error Rate") + theme(legend.position="none") +
  scale_color_manual(name = element_blank(), values = c("line" = "blue", "min" = "red"))
f2_2 = ggplot() +
  geom_line(data = lasso.average.cv.time.and.error, aes(x = fold, y = time, group = 1,
    color = "line")) + ylim(0,65) +
  geom_point(data = lasso.average.cv.time.and.error, aes(x = fold[which.min(time)],
    y = min(time), color = "min")) +
  labs(x = "Cross Validation Folds", y = "Time[sec]", title = "LASSO CV Running Time") +
  theme(legend.position="none") +
  scale_color_manual(name = element_blank(), values = c("line" = "blue", "min" = "red"))
f2 = grid.arrange(f2_1, f2_2, nrow=1, widths = c(1,1))

# cross validation folds for ridge
ridge.average.cv.time.and.error = data.frame(error = round(colMeans(cv_error_nfold[,16:30]),4),
      time = round(colMeans(time.cv.nfold[,16:30]),2),
      fold = factor(c("3","4","5","6","7",
        "8","9","10","15","20",
        "25","30","40","50","L00"),
        levels = c("3","4","5","6","7",
          "8","9","10","15","20",
          "25","30","40","50","L00")))

f3_1 = ggplot() +
  geom_line(data = ridge.average.cv.time.and.error, aes(x = fold, y = error, group = 1,
    color = "line")) + ylim(0.15,.2) +
  geom_point(data = ridge.average.cv.time.and.error, aes(x = fold[which.min(error)],
    y = min(error), color = "min")) +
  labs(x = "Cross Validation Folds", y = "Misclassification Error Rate",
    title = "Ridge CV Error Rate") + theme(legend.position="none") +
  scale_color_manual(name = element_blank(), values = c("line" = "blue", "min" = "red"))
f3_2 = ggplot() +
  geom_line(data = ridge.average.cv.time.and.error, aes(x = fold, y = time, group = 1,
    color = "line")) + ylim(0,65) +
  geom_point(data = ridge.average.cv.time.and.error, aes(x = fold[which.min(time)],
    y = min(time), color = "min")) +
  labs(x = "Cross Validation Folds", y = "Time[sec]", title = "Ridge CV Running Time") +
  theme(legend.position="none") +
  scale_color_manual(name = element_blank(), values = c("line" = "blue", "min" = "red"))
f3 = grid.arrange(f3_1, f3_2, nrow=1, widths = c(1,1))

```

```

# cv error rate with lambda
f4_1 = ggplot() +
  geom_line(data = lasso.cv.error, aes(x = lambda, y = error, group = 1,
                                       color = "line")) + ylim(0,.75) +
  geom_point(data = lasso.cv.error, aes(x = lambda[which.min(error)],
                                       y = min(error), color = "min")) +
  labs(x = "Lambda", y = "CV Misclassification Error Rate",
       title = "LASSO CV Error Rate") + theme(legend.position="none") +
  scale_color_manual(name = element_blank(), values = c("line" = "blue", "min" = "red"))
f4_2 = ggplot() +
  geom_line(data = ridge.cv.error, aes(x = lambda, y = error, group = 1,
                                       color = "line")) + ylim(0,.75) +
  geom_point(data = ridge.cv.error, aes(x = lambda[which.min(error)],
                                       y = min(error), color = "min")) +
  labs(x = "Lambda", y = "CV Misclassification Error Rate",
       title = "Ridge CV Error Rate") + theme(legend.position="none") +
  scale_color_manual(name = element_blank(), values = c("line" = "blue", "min" = "red"))
f4 = grid.arrange(f4_1, f4_2, nrow=1, widths = c(1,1))

#comparison between logistic, tuned lasso, and tuned ridge regression
f5_1 = ggplot(melt(tune_train_error), aes(x = variable, y = value, color = variable)) +
  geom_boxplot() + ylim(0,.5) + labs(x = element_blank(), y = "Misclassification Error Rate",
                                   title = "Train Error Rate with Optimal Lambda") +
  scale_color_brewer(palette="Dark2") + theme(legend.position="none")
f5_2 = ggplot(melt(tune_test_error), aes(x = variable, y = value, color = variable)) +
  geom_boxplot() + ylim(0,.5) + labs(x = element_blank(), y = "Misclassification Error Rate",
                                   title = "Test Error Rate with Optimal Lambda") +
  scale_color_brewer(palette="Dark2") + theme(legend.position="none")
f5 = grid.arrange(f5_1, f5_2, nrow=1, widths = c(1,1))

# variable importance
tune_lasso.coef.mean = data.frame(name = dimnames(coef(lasso.mod)),
                                 coef = colMeans(tune_lasso.coef) )
tune_lasso.coef.mean = tune_lasso.coef.mean[,-2]
tune_ridge.coef.mean = data.frame(name = dimnames(coef(ridge.mod)),
                                 coef = colMeans(tune_ridge.coef) )
tune_ridge.coef.mean = tune_ridge.coef.mean[,-2]

top5.positive.lasso = tune_lasso.coef.mean %>% arrange(-coef) %>% top_n(5)
top5.negative.lasso = tune_lasso.coef.mean %>% arrange(coef) %>% top_n(-5)
top5.lasso = cbind(top5.positive.lasso,top5.negative.lasso)
xtable(top5.lasso)
top5.positive.ridge = tune_ridge.coef.mean %>% arrange(-coef) %>% top_n(5)
top5.negative.ridge = tune_ridge.coef.mean %>% arrange(coef) %>% top_n(-5)
top5.ridge = cbind(top5.positive.ridge,top5.negative.ridge)
xtable(top5.ridge)

f6_1 = ggplot(tune_lasso.coef.mean, aes(x = tune_lasso.coef.mean[,1], y = coef)) +
  geom_bar(stat="identity") + ylim(-1,1) + theme(legend.position="none")+
  labs(x = element_blank(), y = "Average Coefficients",
       title = "LASSO Variable Importance") +
  theme(axis.text.x=element_blank(), axis.ticks.x=element_blank())
f6_2 = ggplot(tune_ridge.coef.mean, aes(x = tune_ridge.coef.mean[,1], y = coef)) +

```

```
geom_bar(stat="identity") + ylim(-1,1) + theme(legend.position="none")+  
labs(x = element_blank(), y = "Average Coefficients",  
     title = "Ridge Variable Importance") +  
theme(axis.text.x=element_blank(), axis.ticks.x=element_blank())  
f6 = grid.arrange(f6_1, f6_2, nrow = 1)
```